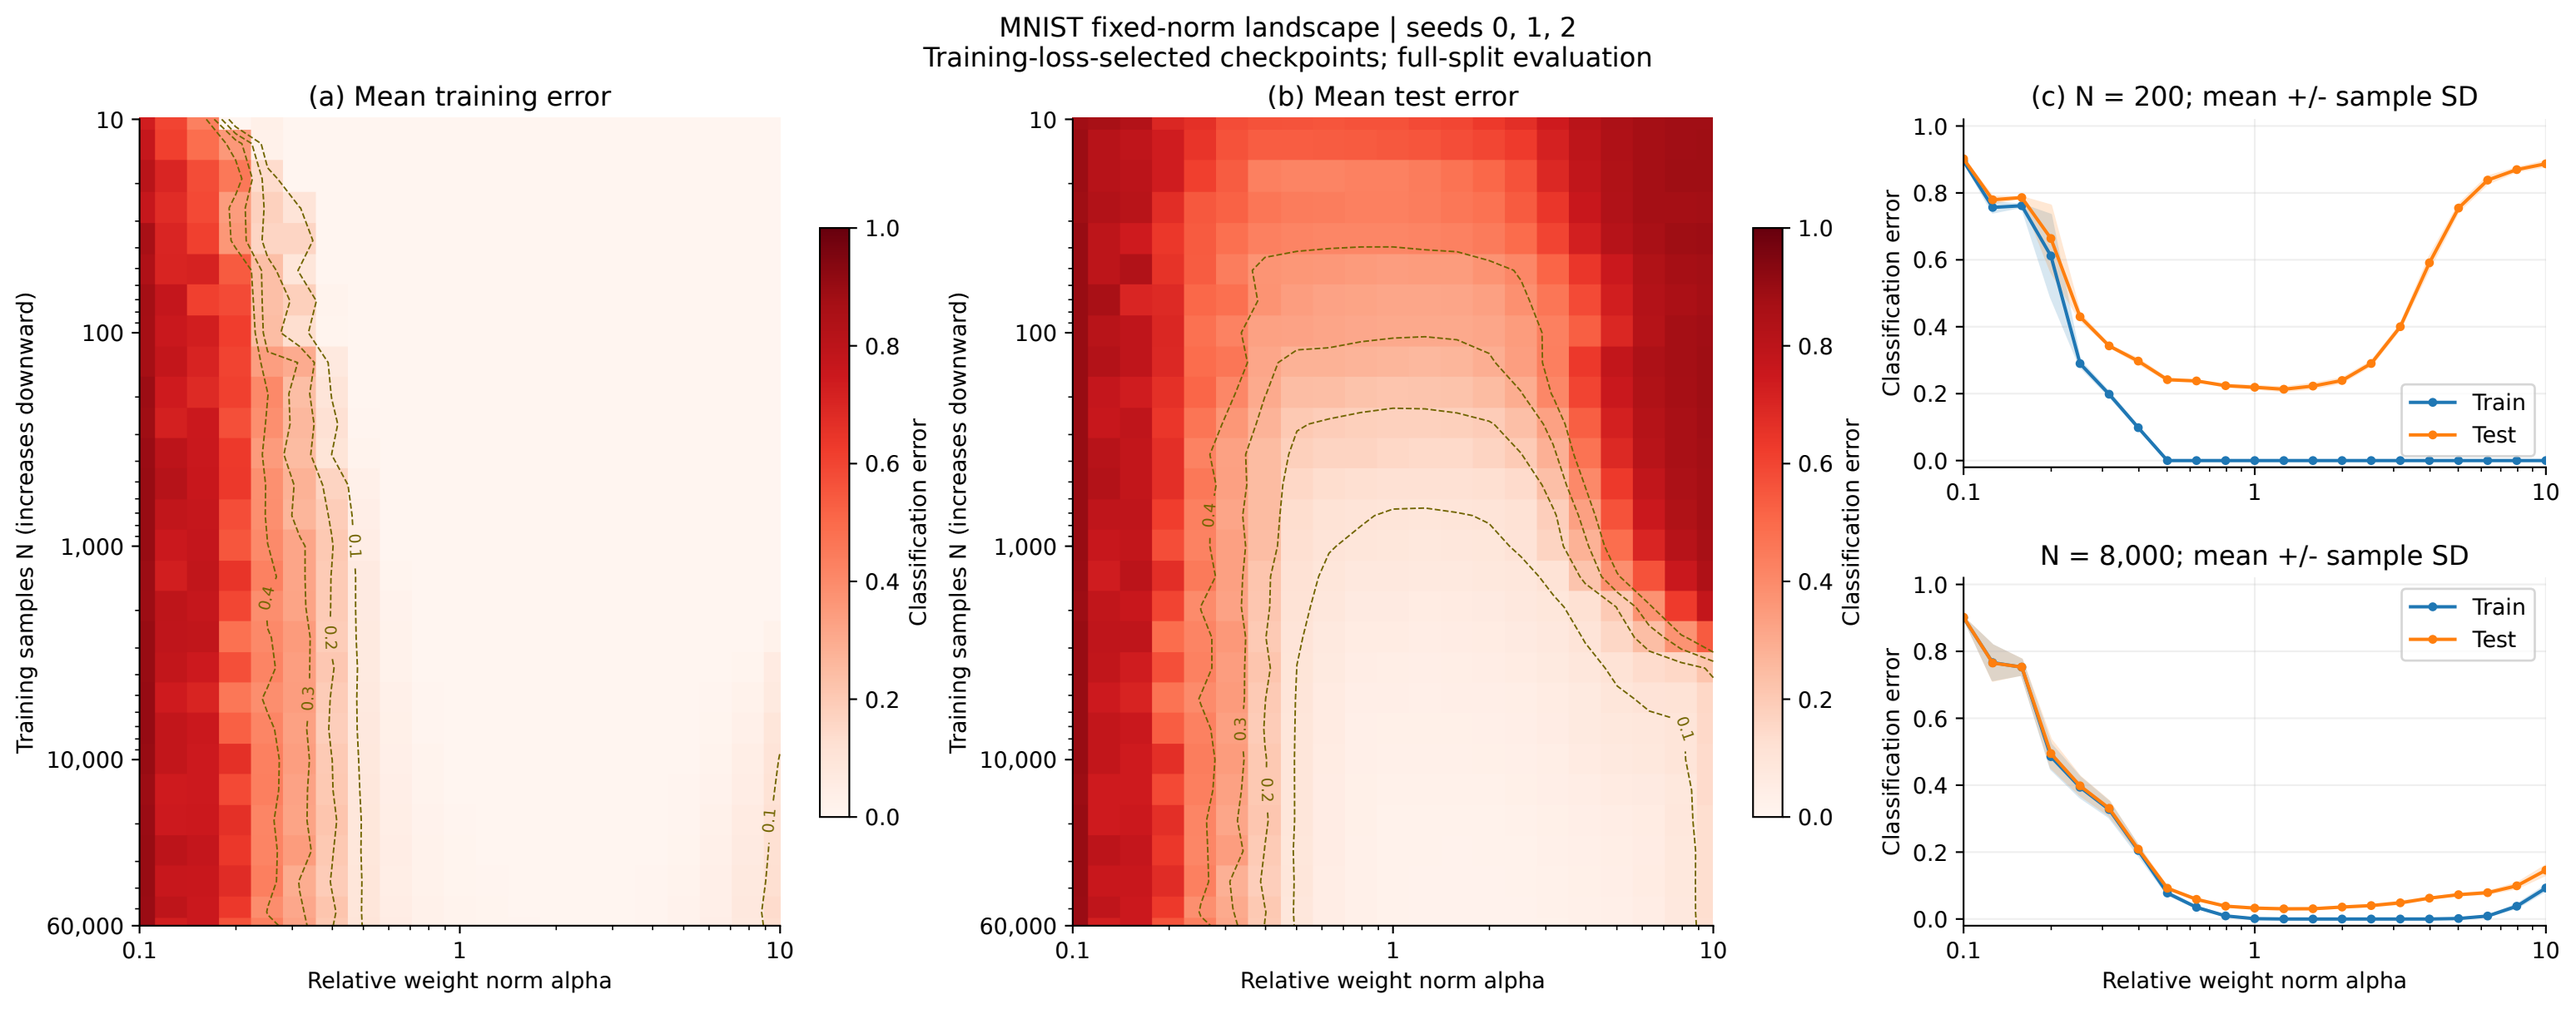
MNIST fixed-norm landscape | seeds 0, 1, 2
Training-loss-selected checkpoints; full-split evaluation
(a) Mean training error
(b) Mean test error
(c) N = 200; mean +/- sample SD
N = 8,000; mean +/- sample SD
Training samples N (increases downward)
Relative weight norm alpha
Classification error
Train
Test
10
100
1,000
10,000
60,000
0.1
1
10
0.0
0.2
0.4
0.6
0.8
1.0
0.1
0.2
0.3
0.4